

```
In [1]: import os
os.environ['KMP_DUPLICATE_LIB_OK']='True'
DATASETS_FOLDER = 'datasets'
if not os.path.isdir(DATASETS_FOLDER):
    os.mkdir(DATASETS_FOLDER)
os.environ['FUEL_DATA_PATH']='./datasets/'
```

```
In [2]: import numpy as np
np.random.seed(7100) # for reproducibility

import warnings
warnings.filterwarnings('ignore',category=FutureWarning)
import tensorflow as tf

from tensorflow import keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, Activation, Flatten
from tensorflow.keras.layers import Conv2D, MaxPooling2D, MaxPooling3D
from tensorflow.keras.constraints import MaxNorm
from tensorflow.keras.optimizers import SGD
from tensorflow.keras import utils

from skimage.transform import resize
import matplotlib.pyplot as plt
from lfw_fuel import lfw
```

Crop and downsample function

```

In [3]: def crop_and_downsample(originalX, downsample_size=32):
        """
        Starts with a 250 x 250 image.
        Crops to 128 x 128 around the center.
        Downsamples the image to (downsample_size) x (downsample_size).
        Returns an image with dimensions (channel, width, height).
        """
        current_dim = 250
        target_dim = 128
        margin = int((current_dim - target_dim)/2)
        left_margin = margin
        right_margin = current_dim - margin

        # newim is shape (6, 128, 128)
        newim = originalX[:, left_margin:right_margin, left_margin:right_m

        # This transpose is mainly useful for plotting with color:
        # Put the images in standard dimension order
        # (width, height, channels)
        sized1 = newim[0:3,:,:]
        sized1 = np.transpose(sized1,(1,2,0))

        sized2 = newim[3:6,:,:]
        sized2 = np.transpose(sized2,(1,2,0))

        # resized are shape (feature_width, feature_height, 3)
        feature_width = feature_height = downsample_size
        resized1 = resize(sized1, (feature_width, feature_height), order=3)
        resized2 = resize(sized2, (feature_width, feature_height), order=3)

        # re-package into a new X entry
        newX = np.concatenate([resized1,resized2], axis=2)

        return newX

a = 0

```

Load the data from LFW web

There are three dataset options: * Original * Funneling * Deep Funneling

We choose Deep Funneling dataset

```

In [4]: # Load the data, shuffled and split between train and test sets
        (X_train_original, y_train_original), (X_test_original, y_test_original)

```

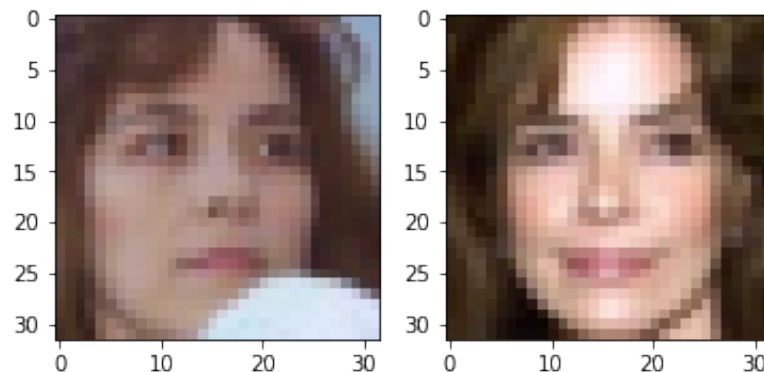
Transform raw pictures

```
In [5]: # Start with 32*32 picture size
ds = 32
X_train = np.asarray([crop_and_downsample(x, downsample_size=ds) for x
X_test  = np.asarray([crop_and_downsample(x, downsample_size=ds) for x
y_train = y_train_original
y_test  = y_test_original
```

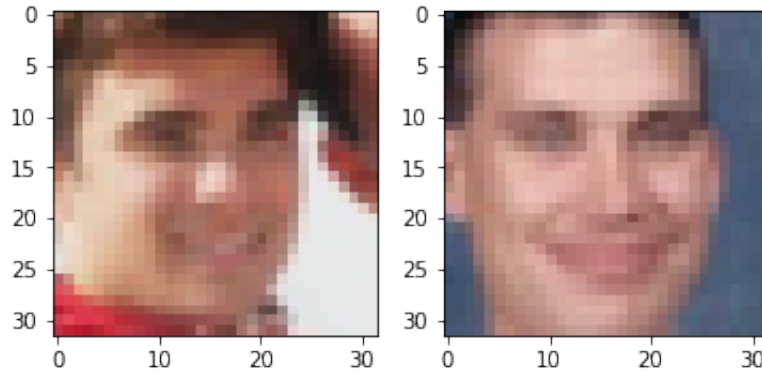
Check Individual Datapoint

Make sure we read the data in correctly

```
In [6]: # datapoint 1
fig = plt.figure()
ax1, ax2 = [fig.add_subplot(1,2,i+1) for i in range(2)]
ax1.imshow(X_train[6,:,:,:0:3])
ax2.imshow(X_train[6,:,:,:3:6])
plt.show()
```



```
In [7]: # datapoint 2
fig = plt.figure()
ax1, ax2 = [fig.add_subplot(1,2,i+1) for i in range(2)]
ax1.imshow(X_test[12,:,:,0:3])
ax2.imshow(X_test[12,:,:,3:6])
plt.show()
```



```
In [8]: print(y_train.shape)
print("Zeros: %d"%(np.sum(y_train==0)))
print("Ones: %d"%(np.sum(y_train==1)))
```

```
(2200, 1)
Zeros: 1100
Ones: 1100
```

Build Baseline Model

- Input
- Conv1: 32 feature, 3*3 kernel size
- Conv2: 64 feature, 3*3 kernel size
- Pooling1: 2*2
- Dropout1: 0.2
- FullConnect
- Dropout2: 0.5
- Output

```
In [9]:
```

Baseline Model

```

base_model = Sequential()

# Input and first convolutional layer
base_model.add(Conv2D(32, (3,3),
                      input_shape=(ds,ds,6),
                      padding='same',
                      data_format='channels_last',
                      activation='relu'))

# Second convolutional layer
base_model.add(Conv2D(64, (3,3),
                      padding='same',
                      data_format='channels_last',
                      activation='relu'))

# Pooling layer 1
base_model.add(AveragePooling2D(pool_size=(2,2),
                                data_format='channels_last'))

# Dropout after pooling 1
base_model.add(Dropout(0.2))

# Flatten layer.
base_model.add(Flatten())

# Fully connected layer
base_model.add(Dense(128, activation='relu', kernel_constraint=MaxNorm(

# Dropout set to 50%.
base_model.add(Dropout(0.5))

# Output layer with 2 units (Y/N) (sigmoid activation function)
base_model.add(Dense(1, activation='sigmoid'))

print(base_model.summary())

```

WARNING:tensorflow:From /Users/nosam/anaconda3/lib/python3.7/site-packages/tensorflow/python/ops/init_ops.py:1251: calling VarianceScaling.__init__ (from tensorflow.python.ops.init_ops) with dtype is deprecated and will be removed in a future version.

Instructions for updating:

Call initializer instance with the dtype argument instead of passing it to the constructor

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	1760

conv2d_1 (Conv2D)	(None, 32, 32, 64)	18496
average_pooling2d (AveragePo	(None, 16, 16, 64)	0
dropout (Dropout)	(None, 16, 16, 64)	0
flatten (Flatten)	(None, 16384)	0
dense (Dense)	(None, 128)	2097280
dropout_1 (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 1)	129
=====		
Total params: 2,117,665		
Trainable params: 2,117,665		
Non-trainable params: 0		
None		

```
In [10]: # Compile model:
base_model.compile(loss='binary_crossentropy', optimizer='adam', metri
```

WARNING:tensorflow:From /Users/nosam/anaconda3/lib/python3.7/site-packages/tensorflow/python/ops/nn_impl.py:180: add_dispatch_support.<locals>.wrapper (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

```
In [11]: epochs = 25
batch_size = 128

base_CNN = base_model.fit(X_train, y_train,
                           batch_size=batch_size,
                           epochs=epochs,
                           verbose=1,
                           validation_data=(X_test, y_test))
```

Train on 2200 samples, validate on 1000 samples

Epoch 1/25

2200/2200 [=====] - 6s 3ms/sample - loss: 0.7164 - binary_accuracy: 0.5086 - val_loss: 0.6863 - val_binary_accuracy: 0.5720

Epoch 2/25

2200/2200 [=====] - 5s 2ms/sample - loss: 0.6756 - binary_accuracy: 0.5818 - val_loss: 0.6515 - val_binary_accuracy: 0.6570

```
Epoch 3/25
2200/2200 [=====] - 5s 2ms/sample - loss: 0.
6496 - binary_accuracy: 0.6305 - val_loss: 0.6285 - val_binary_accu
cy: 0.6490
Epoch 4/25
2200/2200 [=====] - 5s 2ms/sample - loss: 0.
6296 - binary_accuracy: 0.6518 - val_loss: 0.6138 - val_binary_accu
cy: 0.6830
Epoch 5/25
2200/2200 [=====] - 5s 2ms/sample - loss: 0.
6059 - binary_accuracy: 0.6936 - val_loss: 0.6004 - val_binary_accu
cy: 0.6900
Epoch 6/25
2200/2200 [=====] - 5s 2ms/sample - loss: 0.
5991 - binary_accuracy: 0.6941 - val_loss: 0.5978 - val_binary_accu
cy: 0.6860
Epoch 7/25
2200/2200 [=====] - 5s 2ms/sample - loss: 0.
5800 - binary_accuracy: 0.6927 - val_loss: 0.5992 - val_binary_accu
cy: 0.6770
Epoch 8/25
2200/2200 [=====] - 5s 2ms/sample - loss: 0.
5685 - binary_accuracy: 0.7127 - val_loss: 0.5790 - val_binary_accu
cy: 0.6970
Epoch 9/25
2200/2200 [=====] - 5s 2ms/sample - loss: 0.
5398 - binary_accuracy: 0.7305 - val_loss: 0.5830 - val_binary_accu
cy: 0.6990
Epoch 10/25
2200/2200 [=====] - 5s 2ms/sample - loss: 0.
5203 - binary_accuracy: 0.7382 - val_loss: 0.6486 - val_binary_accu
cy: 0.6390
Epoch 11/25
2200/2200 [=====] - 5s 2ms/sample - loss: 0.
5015 - binary_accuracy: 0.7564 - val_loss: 0.6016 - val_binary_accu
cy: 0.6860
Epoch 12/25
2200/2200 [=====] - 5s 2ms/sample - loss: 0.
4928 - binary_accuracy: 0.7709 - val_loss: 0.5980 - val_binary_accu
cy: 0.6810
Epoch 13/25
2200/2200 [=====] - 5s 2ms/sample - loss: 0.
4677 - binary_accuracy: 0.7686 - val_loss: 0.5944 - val_binary_accu
cy: 0.6800
Epoch 14/25
2200/2200 [=====] - 5s 2ms/sample - loss: 0.
4293 - binary_accuracy: 0.7995 - val_loss: 0.6100 - val_binary_accu
cy: 0.6770
Epoch 15/25
2200/2200 [=====] - 5s 2ms/sample - loss: 0.
```

```
4194 - binary_accuracy: 0.8100 - val_loss: 0.6117 - val_binary_accu-  
cy: 0.6840  
Epoch 16/25  
2200/2200 [=====] - 5s 2ms/sample - loss: 0.  
4029 - binary_accuracy: 0.8232 - val_loss: 0.5649 - val_binary_accu-  
cy: 0.7280  
Epoch 17/25  
2200/2200 [=====] - 5s 2ms/sample - loss: 0.  
3744 - binary_accuracy: 0.8395 - val_loss: 0.6114 - val_binary_accu-  
cy: 0.7030  
Epoch 18/25  
2200/2200 [=====] - 5s 2ms/sample - loss: 0.  
3535 - binary_accuracy: 0.8500 - val_loss: 0.5944 - val_binary_accu-  
cy: 0.7000  
Epoch 19/25  
2200/2200 [=====] - 5s 2ms/sample - loss: 0.  
3247 - binary_accuracy: 0.8700 - val_loss: 0.6088 - val_binary_accu-  
cy: 0.7030  
Epoch 20/25  
2200/2200 [=====] - 5s 2ms/sample - loss: 0.  
3127 - binary_accuracy: 0.8659 - val_loss: 0.6125 - val_binary_accu-  
cy: 0.7090  
Epoch 21/25  
2200/2200 [=====] - 5s 2ms/sample - loss: 0.  
2818 - binary_accuracy: 0.8914 - val_loss: 0.6370 - val_binary_accu-  
cy: 0.7080  
Epoch 22/25  
2200/2200 [=====] - 5s 2ms/sample - loss: 0.  
2673 - binary_accuracy: 0.8927 - val_loss: 0.6142 - val_binary_accu-  
cy: 0.7160  
Epoch 23/25  
2200/2200 [=====] - 5s 2ms/sample - loss: 0.  
2576 - binary_accuracy: 0.9023 - val_loss: 0.6491 - val_binary_accu-  
cy: 0.6860  
Epoch 24/25  
2200/2200 [=====] - 5s 2ms/sample - loss: 0.  
2556 - binary_accuracy: 0.9018 - val_loss: 0.7037 - val_binary_accu-  
cy: 0.6950  
Epoch 25/25  
2200/2200 [=====] - 5s 2ms/sample - loss: 0.  
2368 - binary_accuracy: 0.9077 - val_loss: 0.6610 - val_binary_accu-  
cy: 0.7030
```

```
In [12]: res1 = base_model.predict(X_train)  
         res2 = base_model.predict(X_test)
```



```
In [13]: print("Training Data:")
print("Zeros: %d"%(np.sum(res1<0.5)))
print("Ones: %d"%(np.sum(res1>0.5)))
print("\n")
print("Testing Data:")
print("Zeros: %d"%(np.sum(res2<0.5)))
print("Ones: %d"%(np.sum(res2>0.5)))
```

Training Data:
Zeros: 1116
Ones: 1084

Testing Data:
Zeros: 527
Ones: 473


```
In [14]: score = base_model.evaluate(X_test, y_test, verbose=0)
print('Baseline Test accuracy: {0:%}'.format(score[1]))
```

Baseline Test accuracy: 70.300001%

```
In [15]: plt.subplot(3,1,1)
plt.plot(base_CNN.history['binary_accuracy'])
plt.plot(base_CNN.history['val_binary_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='lower right')

plt.subplot(3,1,3)
plt.plot(base_CNN.history['loss'])
plt.plot(base_CNN.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper right')

plt.show()
```



In [16]:

```

modelA = Sequential()

# Input and first convolutional layer
modelA.add(Conv2D(32, (5,5),
                  input_shape=(ds,ds,6),
                  padding='same',
                  data_format='channels_last',
                  activation='relu'))

# Second convolutional layer
modelA.add(Conv2D(64, (5,5),
                  padding='same',
                  data_format='channels_last',
                  activation='relu'))

# Pooling layer 1
modelA.add(AveragePooling2D(pool_size=(2,2),
                             data_format='channels_last'))

# Dropout after pooling 1
modelA.add(Dropout(0.2))

# Flatten layer.
modelA.add(Flatten())

# Fully connected layer
modelA.add(Dense(128, activation='relu', kernel_constraint=MaxNorm(3)))

# Dropout set to 50%.
modelA.add(Dropout(0.5))

# Output layer with 2 units (Y/N) (sigmoid activation function)
modelA.add(Dense(1, activation='sigmoid'))

print(modelA.summary())

```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_2 (Conv2D)	(None, 32, 32, 32)	4832
conv2d_3 (Conv2D)	(None, 32, 32, 64)	51264
average_pooling2d_1 (Average	(None, 16, 16, 64)	0
dropout_2 (Dropout)	(None, 16, 16, 64)	0
flatten_1 (Flatten)	(None, 16384)	0

dense_2 (Dense)	(None, 128)	2097280
dropout_3 (Dropout)	(None, 128)	0
dense_3 (Dense)	(None, 1)	129
=====		
Total params: 2,153,505		
Trainable params: 2,153,505		
Non-trainable params: 0		
None		

```
In [17]: # Compile model:
modelA.compile(loss='binary_crossentropy', optimizer='adam', metrics=[
```

```
In [18]: epochs = 25
batch_size = 128

modelA_CNN = modelA.fit(X_train, y_train,
                        batch_size=batch_size,
                        epochs=epochs,
                        verbose=1,
                        validation_data=(X_test, y_test))
```

Train on 2200 samples, validate on 1000 samples

Epoch 1/25

2200/2200 [=====] - 9s 4ms/sample - loss: 0.7250 - binary_accuracy: 0.5159 - val_loss: 0.6944 - val_binary_accuracy: 0.5000

Epoch 2/25

2200/2200 [=====] - 8s 4ms/sample - loss: 0.6946 - binary_accuracy: 0.5036 - val_loss: 0.6929 - val_binary_accuracy: 0.5430

Epoch 3/25

2200/2200 [=====] - 8s 4ms/sample - loss: 0.6933 - binary_accuracy: 0.5050 - val_loss: 0.6907 - val_binary_accuracy: 0.5030

Epoch 4/25

2200/2200 [=====] - 8s 4ms/sample - loss: 0.6866 - binary_accuracy: 0.5291 - val_loss: 0.6830 - val_binary_accuracy: 0.5660

Epoch 5/25

2200/2200 [=====] - 8s 4ms/sample - loss: 0.6746 - binary_accuracy: 0.5768 - val_loss: 0.6719 - val_binary_accuracy: 0.5820

Epoch 6/25

2200/2200 [=====] - 8s 4ms/sample - loss: 0.6715 - binary_accuracy: 0.5818 - val_loss: 0.6725 - val_binary_accuracy: 0.5690

```
Epoch 7/25
2200/2200 [=====] - 8s 4ms/sample - loss: 0.
6671 - binary_accuracy: 0.5768 - val_loss: 0.6654 - val_binary_accu
cy: 0.5800
Epoch 8/25
2200/2200 [=====] - 8s 4ms/sample - loss: 0.
6495 - binary_accuracy: 0.6141 - val_loss: 0.6568 - val_binary_accu
cy: 0.6290
Epoch 9/25
2200/2200 [=====] - 8s 4ms/sample - loss: 0.
6395 - binary_accuracy: 0.6336 - val_loss: 0.6280 - val_binary_accu
cy: 0.6350
Epoch 10/25
2200/2200 [=====] - 8s 4ms/sample - loss: 0.
6274 - binary_accuracy: 0.6564 - val_loss: 0.5988 - val_binary_accu
cy: 0.6850
Epoch 11/25
2200/2200 [=====] - 9s 4ms/sample - loss: 0.
6009 - binary_accuracy: 0.6868 - val_loss: 0.5995 - val_binary_accu
cy: 0.6800
Epoch 12/25
2200/2200 [=====] - 9s 4ms/sample - loss: 0.
6032 - binary_accuracy: 0.6859 - val_loss: 0.6024 - val_binary_accu
cy: 0.6830
Epoch 13/25
2200/2200 [=====] - 9s 4ms/sample - loss: 0.
5958 - binary_accuracy: 0.6900 - val_loss: 0.5877 - val_binary_accu
cy: 0.6910
Epoch 14/25
2200/2200 [=====] - 9s 4ms/sample - loss: 0.
5535 - binary_accuracy: 0.7277 - val_loss: 0.5673 - val_binary_accu
cy: 0.7010
Epoch 15/25
2200/2200 [=====] - 10s 4ms/sample - loss: 0
.5464 - binary_accuracy: 0.7305 - val_loss: 0.5782 - val_binary_accu
acy: 0.6960
Epoch 16/25
2200/2200 [=====] - 11s 5ms/sample - loss: 0
.5135 - binary_accuracy: 0.7509 - val_loss: 0.5612 - val_binary_accu
acy: 0.7060
Epoch 17/25
2200/2200 [=====] - 10s 5ms/sample - loss: 0
.5156 - binary_accuracy: 0.7459 - val_loss: 0.5770 - val_binary_accu
acy: 0.6930
Epoch 18/25
2200/2200 [=====] - 10s 5ms/sample - loss: 0
.4984 - binary_accuracy: 0.7609 - val_loss: 0.5712 - val_binary_accu
acy: 0.7020
Epoch 19/25
2200/2200 [=====] - 10s 5ms/sample - loss: 0
```

```
.4618 - binary_accuracy: 0.7805 - val_loss: 0.5690 - val_binary_accu
acy: 0.7040
Epoch 20/25
2200/2200 [=====] - 10s 5ms/sample - loss: 0
.4449 - binary_accuracy: 0.7864 - val_loss: 0.5695 - val_binary_accu
acy: 0.7150
Epoch 21/25
2200/2200 [=====] - 11s 5ms/sample - loss: 0
.4272 - binary_accuracy: 0.7986 - val_loss: 0.6437 - val_binary_accu
acy: 0.6720
Epoch 22/25
2200/2200 [=====] - 10s 5ms/sample - loss: 0
.4170 - binary_accuracy: 0.8095 - val_loss: 0.6158 - val_binary_accu
acy: 0.6860
Epoch 23/25
2200/2200 [=====] - 11s 5ms/sample - loss: 0
.4038 - binary_accuracy: 0.8150 - val_loss: 0.6130 - val_binary_accu
acy: 0.6950
Epoch 24/25
2200/2200 [=====] - 11s 5ms/sample - loss: 0
.3704 - binary_accuracy: 0.8305 - val_loss: 0.6330 - val_binary_accu
acy: 0.6820
Epoch 25/25
2200/2200 [=====] - 10s 5ms/sample - loss: 0
.3385 - binary_accuracy: 0.8473 - val_loss: 0.6474 - val_binary_accu
acy: 0.7000
```

```
In [19]: res1 = modelA.predict(X_train)
         res2 = modelA.predict(X_test)
```

```
In [20]: print("Training Data:")
         print("Zeros: %d"%(np.sum(res1<0.5)))
         print("Ones: %d"%(np.sum(res1>0.5)))
         print("\n")
         print("Testing Data:")
         print("Zeros: %d"%(np.sum(res2<0.5)))
         print("Ones: %d"%(np.sum(res2>0.5)))
```

```
Training Data:
Zeros: 1095
Ones: 1105
```

```
Testing Data:
Zeros: 534
Ones: 466
```


```
In [21]: score = modelA.evaluate(X_test, y_test, verbose=0)
print('modelA Test accuracy: {0:%}'.format(score[1]))
```

modelA Test accuracy: 69.999999%

```
In [22]: plt.subplot(3,1,1)
plt.plot(modelA_CNN.history['binary_accuracy'])
plt.plot(modelA_CNN.history['val_binary_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='lower right')

plt.subplot(3,1,3)
plt.plot(modelA_CNN.history['loss'])
plt.plot(modelA_CNN.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper right')

plt.show()
```



```
In [23]:
```

```

modelB = Sequential()

# Input and first convolutional layer
modelB.add(Conv2D(32, (7,7),
                  input_shape=(ds,ds,6),
                  padding='same',
                  data_format='channels_last',
                  activation='relu'))

# Second convolutional layer
modelB.add(Conv2D(64, (7,7),
                  padding='same',
                  data_format='channels_last',
                  activation='relu'))

# Pooling layer 1
modelB.add(AveragePooling2D(pool_size=(2,2),
                             data_format='channels_last'))

# Dropout after pooling 1
modelB.add(Dropout(0.2))

# Flatten layer.
modelB.add(Flatten())

# Fully connected layer
modelB.add(Dense(128, activation='relu', kernel_constraint=MaxNorm(3)))

# Dropout set to 50%.
modelB.add(Dropout(0.5))

# Output layer with 2 units (Y/N) (sigmoid activation function)
modelB.add(Dense(1, activation='sigmoid'))

print(modelB.summary())

```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
conv2d_4 (Conv2D)	(None, 32, 32, 32)	9440
conv2d_5 (Conv2D)	(None, 32, 32, 64)	100416
average_pooling2d_2 (Average	(None, 16, 16, 64)	0
dropout_4 (Dropout)	(None, 16, 16, 64)	0
flatten_2 (Flatten)	(None, 16384)	0

dense_4 (Dense)	(None, 128)	2097280
dropout_5 (Dropout)	(None, 128)	0
dense_5 (Dense)	(None, 1)	129
=====		
Total params: 2,207,265		
Trainable params: 2,207,265		
Non-trainable params: 0		
None		

```
In [24]: # Compile model:
modelB.compile(loss='binary_crossentropy', optimizer='adam', metrics=[
```

```
In [25]: epochs = 25
batch_size = 128

modelB_CNN = modelB.fit(X_train, y_train,
                        batch_size=batch_size,
                        epochs=epochs,
                        verbose=1,
                        validation_data=(X_test, y_test))
```

Train on 2200 samples, validate on 1000 samples

Epoch 1/25

2200/2200 [=====] - 15s 7ms/sample - loss: 0.7179 - binary_accuracy: 0.4950 - val_loss: 0.6895 - val_binary_accuracy: 0.5620

Epoch 2/25

2200/2200 [=====] - 14s 6ms/sample - loss: 0.6823 - binary_accuracy: 0.5505 - val_loss: 0.6764 - val_binary_accuracy: 0.5580

Epoch 3/25

2200/2200 [=====] - 14s 6ms/sample - loss: 0.6852 - binary_accuracy: 0.5482 - val_loss: 0.7093 - val_binary_accuracy: 0.5360

Epoch 4/25

2200/2200 [=====] - 14s 6ms/sample - loss: 0.6724 - binary_accuracy: 0.5895 - val_loss: 0.6439 - val_binary_accuracy: 0.6620

Epoch 5/25

2200/2200 [=====] - 14s 6ms/sample - loss: 0.6588 - binary_accuracy: 0.6245 - val_loss: 0.6374 - val_binary_accuracy: 0.6670

Epoch 6/25

2200/2200 [=====] - 14s 6ms/sample - loss: 0.6492 - binary_accuracy: 0.6259 - val_loss: 0.6187 - val_binary_accuracy: 0.6490

```
Epoch 7/25
2200/2200 [=====] - 15s 7ms/sample - loss: 0
.6272 - binary_accuracy: 0.6559 - val_loss: 0.6094 - val_binary_accu
acy: 0.6790
Epoch 8/25
2200/2200 [=====] - 16s 7ms/sample - loss: 0
.6240 - binary_accuracy: 0.6586 - val_loss: 0.6066 - val_binary_accu
acy: 0.6800
Epoch 9/25
2200/2200 [=====] - 17s 8ms/sample - loss: 0
.6242 - binary_accuracy: 0.6505 - val_loss: 0.6441 - val_binary_accu
acy: 0.6120
Epoch 10/25
2200/2200 [=====] - 19s 8ms/sample - loss: 0
.5980 - binary_accuracy: 0.6855 - val_loss: 0.5825 - val_binary_accu
acy: 0.7050
Epoch 11/25
2200/2200 [=====] - 20s 9ms/sample - loss: 0
.5864 - binary_accuracy: 0.6950 - val_loss: 0.5843 - val_binary_accu
acy: 0.7140
Epoch 12/25
2200/2200 [=====] - 19s 9ms/sample - loss: 0
.5659 - binary_accuracy: 0.7141 - val_loss: 0.6659 - val_binary_accu
acy: 0.6380
Epoch 13/25
2200/2200 [=====] - 20s 9ms/sample - loss: 0
.5691 - binary_accuracy: 0.7095 - val_loss: 0.5713 - val_binary_accu
acy: 0.6970
Epoch 14/25
2200/2200 [=====] - 19s 8ms/sample - loss: 0
.5441 - binary_accuracy: 0.7245 - val_loss: 0.5677 - val_binary_accu
acy: 0.7150
Epoch 15/25
2200/2200 [=====] - 18s 8ms/sample - loss: 0
.5206 - binary_accuracy: 0.7514 - val_loss: 0.6433 - val_binary_accu
acy: 0.6540
Epoch 16/25
2200/2200 [=====] - 17s 8ms/sample - loss: 0
.5147 - binary_accuracy: 0.7509 - val_loss: 0.5686 - val_binary_accu
acy: 0.7200
Epoch 17/25
2200/2200 [=====] - 17s 8ms/sample - loss: 0
.4686 - binary_accuracy: 0.7723 - val_loss: 0.5632 - val_binary_accu
acy: 0.7090
Epoch 18/25
2200/2200 [=====] - 17s 8ms/sample - loss: 0
.4446 - binary_accuracy: 0.7941 - val_loss: 0.5719 - val_binary_accu
acy: 0.7130
Epoch 19/25
2200/2200 [=====] - 17s 8ms/sample - loss: 0
```

```
.4140 - binary_accuracy: 0.8086 - val_loss: 0.5621 - val_binary_accu
acy: 0.7220
Epoch 20/25
2200/2200 [=====] - 18s 8ms/sample - loss: 0
.4133 - binary_accuracy: 0.8182 - val_loss: 0.5941 - val_binary_accu
acy: 0.7090
Epoch 21/25
2200/2200 [=====] - 18s 8ms/sample - loss: 0
.3986 - binary_accuracy: 0.8227 - val_loss: 0.5830 - val_binary_accu
acy: 0.7300
Epoch 22/25
2200/2200 [=====] - 18s 8ms/sample - loss: 0
.3567 - binary_accuracy: 0.8414 - val_loss: 0.6037 - val_binary_accu
acy: 0.7130
Epoch 23/25
2200/2200 [=====] - 18s 8ms/sample - loss: 0
.3099 - binary_accuracy: 0.8705 - val_loss: 0.6107 - val_binary_accu
acy: 0.7160
Epoch 24/25
2200/2200 [=====] - 18s 8ms/sample - loss: 0
.2754 - binary_accuracy: 0.8832 - val_loss: 0.6547 - val_binary_accu
acy: 0.7060
Epoch 25/25
2200/2200 [=====] - 18s 8ms/sample - loss: 0
.2590 - binary_accuracy: 0.8918 - val_loss: 0.6686 - val_binary_accu
acy: 0.7040
```

```
In [26]: res1 = modelB.predict(X_train)
         res2 = modelB.predict(X_test)
```

```
In [27]: print("Training Data:")
         print("Zeros: %d"%(np.sum(res1<0.5)))
         print("Ones: %d"%(np.sum(res1>0.5)))
         print("\n")
         print("Testing Data:")
         print("Zeros: %d"%(np.sum(res2<0.5)))
         print("Ones: %d"%(np.sum(res2>0.5)))
```

```
Training Data:
Zeros: 1167
Ones: 1033
```

```
Testing Data:
Zeros: 598
Ones: 402
```


```
In [28]: score = modelB.evaluate(X_test, y_test, verbose=0)
print('modelB Test accuracy: {0:%}'.format(score[1]))
```

modelB Test accuracy: 70.400000%

```
In [29]: plt.subplot(3,1,1)
plt.plot(modelB_CNN.history['binary_accuracy'])
plt.plot(modelB_CNN.history['val_binary_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='lower right')

plt.subplot(3,1,3)
plt.plot(modelB_CNN.history['loss'])
plt.plot(modelB_CNN.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper right')

plt.show()
```



Build More Complex Model

Compare baseline, modelA and modelB, we find that overall modelA is the best model(less underfitting). So we build more complex model on top of modelA

Model C:

- Input
- Conv1: 32 feature, 5*5 kernel size
- Conv2: 64 feature, 5*5 kernel size
- Pooling1: 2*2
- Dropout1: 0.2
- Conv3: 128 feature, 5*5 kernel size
- Pooling2: 2*2
- Dropout2: 0.2
- FullConnect
- Dropout3: 0.5
- Output

```
In [30]: # modelC

modelC = Sequential()

# Input and first convolutional layer
modelC.add(Conv2D(32, (5,5),
                  input_shape=(ds,ds,6),
                  padding='same',
                  data_format='channels_last',
                  activation='relu'))

# Second convolutional layer
modelC.add(Conv2D(64, (5,5),
                  padding='same',
                  data_format='channels_last',
                  activation='relu'))

# Pooling layer 1
modelC.add(AveragePooling2D(pool_size=(2,2),
                             data_format='channels_last'))

# Dropout after pooling 1
modelC.add(Dropout(0.2))

# Third convolutional layer
modelC.add(Conv2D(128, (5,5),
                  padding='same',
```

```

padding='same',
data_format='channels_last',
activation='relu'))

# Pooling layer 2
modelC.add(AveragePooling2D(pool_size=(2,2),
                             data_format='channels_last'))

# Dropout after pooling 2
modelC.add(Dropout(0.2))

# Flatten layer.
modelC.add(Flatten())

# Fully connected layer
modelC.add(Dense(128, activation='relu', kernel_constraint=MaxNorm(3)))

# Dropout set to 50%.
modelC.add(Dropout(0.5))

# Output layer with 2 units (Y/N) (sigmoid activation function)
modelC.add(Dense(1, activation='sigmoid'))

print(modelC.summary())

```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
conv2d_6 (Conv2D)	(None, 32, 32, 32)	4832
conv2d_7 (Conv2D)	(None, 32, 32, 64)	51264
average_pooling2d_3 (Average	(None, 16, 16, 64)	0
dropout_6 (Dropout)	(None, 16, 16, 64)	0
conv2d_8 (Conv2D)	(None, 16, 16, 128)	204928
average_pooling2d_4 (Average	(None, 8, 8, 128)	0
dropout_7 (Dropout)	(None, 8, 8, 128)	0
flatten_3 (Flatten)	(None, 8192)	0
dense_6 (Dense)	(None, 128)	1048704
dropout_8 (Dropout)	(None, 128)	0
dense_7 (Dense)	(None, 1)	129

Total params: 1,309,857
 Trainable params: 1,309,857
 Non-trainable params: 0

None

```
In [31]: # Compile model:
modelC.compile(loss='binary_crossentropy', optimizer='adam', metrics=[
```

```
In [32]: epochs = 25
batch_size = 128

modelC_CNN = modelC.fit(X_train, y_train,
                        batch_size=batch_size,
                        epochs=epochs,
                        verbose=1,
                        validation_data=(X_test, y_test))
```

Train on 2200 samples, validate on 1000 samples

Epoch 1/25

2200/2200 [=====] - 17s 8ms/sample - loss: 0.6990 - binary_accuracy: 0.4818 - val_loss: 0.6931 - val_binary_accuracy: 0.5000

Epoch 2/25

2200/2200 [=====] - 16s 7ms/sample - loss: 0.6932 - binary_accuracy: 0.4977 - val_loss: 0.6928 - val_binary_accuracy: 0.5060

Epoch 3/25

2200/2200 [=====] - 15s 7ms/sample - loss: 0.6894 - binary_accuracy: 0.5350 - val_loss: 0.6872 - val_binary_accuracy: 0.5400

Epoch 4/25

2200/2200 [=====] - 15s 7ms/sample - loss: 0.6878 - binary_accuracy: 0.5500 - val_loss: 0.6687 - val_binary_accuracy: 0.6230

Epoch 5/25

2200/2200 [=====] - 15s 7ms/sample - loss: 0.6643 - binary_accuracy: 0.5941 - val_loss: 0.6386 - val_binary_accuracy: 0.6390

Epoch 6/25

2200/2200 [=====] - 15s 7ms/sample - loss: 0.6488 - binary_accuracy: 0.6336 - val_loss: 0.6237 - val_binary_accuracy: 0.6560

Epoch 7/25

2200/2200 [=====] - 16s 7ms/sample - loss: 0.6421 - binary_accuracy: 0.6295 - val_loss: 0.6169 - val_binary_accuracy: 0.6860

Epoch 8/25

2200/2200 [=====] - 18s 8ms/sample - loss: 0

```
.6435 - binary_accuracy: 0.6309 - val_loss: 0.6115 - val_binary_accu
acy: 0.6800
Epoch 9/25
2200/2200 [=====] - 21s 9ms/sample - loss: 0
.6139 - binary_accuracy: 0.6573 - val_loss: 0.6155 - val_binary_accu
acy: 0.6670
Epoch 10/25
2200/2200 [=====] - 23s 11ms/sample - loss:
0.5989 - binary_accuracy: 0.6845 - val_loss: 0.5847 - val_binary_accu
racy: 0.6890
Epoch 11/25
2200/2200 [=====] - 24s 11ms/sample - loss:
0.5859 - binary_accuracy: 0.6945 - val_loss: 0.5723 - val_binary_accu
racy: 0.7070
Epoch 12/25
2200/2200 [=====] - 21s 9ms/sample - loss: 0
.5847 - binary_accuracy: 0.6918 - val_loss: 0.5877 - val_binary_accu
acy: 0.6990
Epoch 13/25
2200/2200 [=====] - 18s 8ms/sample - loss: 0
.5602 - binary_accuracy: 0.7173 - val_loss: 0.5627 - val_binary_accu
acy: 0.7180
Epoch 14/25
2200/2200 [=====] - 17s 8ms/sample - loss: 0
.5605 - binary_accuracy: 0.7159 - val_loss: 0.5628 - val_binary_accu
acy: 0.7030
Epoch 15/25
2200/2200 [=====] - 17s 8ms/sample - loss: 0
.5309 - binary_accuracy: 0.7286 - val_loss: 0.5547 - val_binary_accu
acy: 0.7210
Epoch 16/25
2200/2200 [=====] - 17s 8ms/sample - loss: 0
.5244 - binary_accuracy: 0.7286 - val_loss: 0.5815 - val_binary_accu
acy: 0.7010
Epoch 17/25
2200/2200 [=====] - 17s 8ms/sample - loss: 0
.5037 - binary_accuracy: 0.7523 - val_loss: 0.6131 - val_binary_accu
acy: 0.6810
Epoch 18/25
2200/2200 [=====] - 17s 8ms/sample - loss: 0
.4909 - binary_accuracy: 0.7564 - val_loss: 0.5689 - val_binary_accu
acy: 0.7130
Epoch 19/25
2200/2200 [=====] - 17s 8ms/sample - loss: 0
.4914 - binary_accuracy: 0.7605 - val_loss: 0.5551 - val_binary_accu
acy: 0.7170
Epoch 20/25
2200/2200 [=====] - 17s 8ms/sample - loss: 0
.4599 - binary_accuracy: 0.7795 - val_loss: 0.6063 - val_binary_accu
acy: 0.7020
```



```

Epoch 21/25
2200/2200 [=====] - 22s 10ms/sample - loss: 0.4515 - binary_accuracy: 0.7882 - val_loss: 0.6424 - val_binary_accuracy: 0.6810
Epoch 22/25
2200/2200 [=====] - 19s 9ms/sample - loss: 0.4518 - binary_accuracy: 0.7918 - val_loss: 0.6162 - val_binary_accuracy: 0.7240
Epoch 23/25
2200/2200 [=====] - 18s 8ms/sample - loss: 0.4136 - binary_accuracy: 0.8041 - val_loss: 0.6017 - val_binary_accuracy: 0.7150
Epoch 24/25
2200/2200 [=====] - 18s 8ms/sample - loss: 0.3993 - binary_accuracy: 0.8150 - val_loss: 0.6171 - val_binary_accuracy: 0.7080
Epoch 25/25
2200/2200 [=====] - 17s 8ms/sample - loss: 0.4171 - binary_accuracy: 0.8023 - val_loss: 0.6260 - val_binary_accuracy: 0.7000

```

```

In [33]: res1 = modelC.predict(X_train)
         res2 = modelC.predict(X_test)

```

```

In [34]: print("Training Data:")
         print("Zeros: %d"%(np.sum(res1<0.5)))
         print("Ones: %d"%(np.sum(res1>0.5)))
         print("\n")
         print("Testing Data:")
         print("Zeros: %d"%(np.sum(res2<0.5)))
         print("Ones: %d"%(np.sum(res2>0.5)))

```

```

Training Data:
Zeros: 1235
Ones: 965

```

```

Testing Data:
Zeros: 584
Ones: 416

```

```

In [35]: score = modelC.evaluate(X_test, y_test, verbose=0)
         print('modelC Test accuracy: {0:%}'.format(score[1]))


modelC Test accuracy: 69.999999%

```

```
In [36]: plt.subplot(3,1,1)
plt.plot(modelC_CNN.history['binary_accuracy'])
plt.plot(modelC_CNN.history['val_binary_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='lower right')

plt.subplot(3,1,3)
plt.plot(modelC_CNN.history['loss'])
plt.plot(modelC_CNN.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper right')

plt.show()
```



Final Best Model

We experiment batch_size value: 32, 64, 128, 256 And we find, batch_size=128 is the best
So we use callback function to find the best model on top of modelC

```
In [37]: from tensorflow.keras.callbacks import ModelCheckpoint
# Define a checkpoint to save the data
checkpoint_name = 'weights.best.hdf5'
checkpoint = ModelCheckpoint(checkpoint_name, monitor='val_binary_accu
callbacks_list = [checkpoint]
```

```
In [38]: # Train the model
final_CNN = modelC.fit(X_train, y_train,
                        batch_size=128,
                        epochs=100,
                        verbose=1,
                        validation_data=(X_test, y_test),
                        callbacks=callbacks_list)
```

Train on 2200 samples, validate on 1000 samples

Epoch 1/100

2176/2200 [=====>.] - ETA: 0s - loss: 0.3833 - binary_accuracy: 0.8258

Epoch 00001: val_binary_accuracy improved from -inf to 0.69400, saving model to weights.best.hdf5

2200/2200 [=====] - 15s 7ms/sample - loss: 0.3838 - binary_accuracy: 0.8250 - val_loss: 0.6055 - val_binary_accuracy: 0.6940

Epoch 2/100

2176/2200 [=====>.] - ETA: 0s - loss: 0.3529 - binary_accuracy: 0.8415

Epoch 00002: val_binary_accuracy improved from 0.69400 to 0.71200, saving model to weights.best.hdf5

2200/2200 [=====] - 15s 7ms/sample - loss: 0.3522 - binary_accuracy: 0.8418 - val_loss: 0.6078 - val_binary_accuracy: 0.7120

Epoch 3/100

2176/2200 [=====>.] - ETA: 0s - loss: 0.3178 - binary_accuracy: 0.8566

```
In [39]: # Load wights file of the best model :
weights_file = 'weights.best.hdf5'
best_CNN = modelC.load_weights(weights_file) # load it
modelC.compile(loss='binary_crossentropy', optimizer='adam', metrics=[
```

```
In [40]: score = modelC.evaluate(X_test, y_test, verbose=0)
print('Best CNN Model Test accuracy: {0:%}'.format(score[1]))
```

Best CNN Model Test accuracy: 72.799999%

```
In [41]: # Make predictions
predictions = modelC.predict(X_test)
```

In [42]: predictions


```
Out[42]: array([[3.36706936e-01],
                [6.15119934e-05],
                [9.56135035e-01],
                [2.48032808e-03],
                [1.13038272e-01],
                [9.98085678e-01],
                [3.25866759e-01],
                [9.54843640e-01],
                [1.00000000e+00],
                [8.95335317e-01],
                [3.47958922e-01],
                [2.16363668e-02],
                [1.44955635e-01],
                [9.99999881e-01],
                [2.69778967e-02],
                [7.69262195e-01],
                [9.07015562e-01],
                [9.99962389e-01],
                [5.96199572e-01],
                [0.00000000e+00]])
```

```
In [43]: plt.style.use('ggplot')

def plot_history(history):
    loss = history.history['loss']
    val_loss = history.history['val_loss']
    x = range(1, len(loss) + 1)

    plt.figure(figsize=(12, 5))
    plt.plot(x, loss, 'b', label='Training loss')
    plt.plot(x, val_loss, 'r', label='Validation loss')
    plt.title('Training and validation loss')
    plt.legend()
    plt.show()

plot_history(final_CNN)
```



End